

CSE 145 - Homework 2

Daniel Xiong (dxiong5@ucsc.edu)

Due May 14

1 Introduction

The goal of this project was to data mine the `Customer_Churn.xlsx` dataset, which contains 20,000 entries with 12 variables describing features of customers of a mobile phone provider. We aimed to predict the variable `Leave`, which represented whether a given customer would stay or leave the company. To achieve this goal, we first create some visualizations to try and understand the data. We then used the k -means algorithm to cluster the data to further analyze the properties of the data. Finally, we chose predictive models to predict whether or not a customer would stay or leave the company.

2 Tools Used

This assignment was completed using Python 3.7 (`sklearn`, `keras`, `pandas`, `matplotlib`).

3 Data Pre-processing

The `Customer_Churn.xlsx` spreadsheet had attribute values that were strings, such as "LEAVE" and "STAY" for the `Leave` attribute. To convert these into numerical values, I used label encoding. The specific encodings can be found in `../data/label_encodings.txt`.

4 Data Visualization

I decided to visualize information gain, which is a measure of how much an attribute improves entropy over the whole segmentation it creates. **Figure 1** is a bar graph that ranks the information gain of all the attributes in decreasing order, which would help us understand which attributes are the most important.

The next visualization I decided to create is a heatmap that shows correlations between attributes. Shown in **Figure 2**, the heatmap is useful for understanding how certain attributes are related to each other, and can help understand how to increase customer churn.

The code for these visualizations can be found in `visualizations.py`.




Figure 1: Attributes ranked by their information gain




Figure 2: Correlation Heatmap

5 Customer Segmentation with k -means

Before training a k -means model on the dataset, I first had to scale the data so that the magnitudes would not be so vastly different. To do this, I used the `StandardScaler` function from Python's `sklearn` module.

I then used `sklearn`'s `KMeans` function for the k -means model along with the `k-means++` centroid initializer.

In order to determine the best value of k , I trained many different k -means models each with a different k from $k = 2 \rightarrow 25$. I then created two plots: a Sum Squared Error (SSE) plot and a Silhouette plot, **Figure 3** and **Figure 4**, respectively. I used the elbow method, along with the silhouette plot, to determine that a k -means model with $k = 5$ would result in the optimal clustering.

The code for the k -means model and its analysis can be found in `../src/kmeans.py`.




Figure 3: SSE plot




Figure 4: Silhouette plot

5.1 Cluster Analysis ($k = 5$)

The following tables contain the mean and standard deviation for each attribute and number of points in each of the 5 clusters. The raw data, which contains more information in the form of a five-number summary for each attribute (I did not include this due to space), can be found in `../data/cluster_descriptions.txt`. One thing that I would like to note is that since many of the attributes were label encoded, the mean for such attributes would represent the average frequency of attribute values in the cluster.

From the tables, we can see that the attributes **College**, **House**, **Reported Satisfaction**, **Reported Usage Level**, and **Considering Change of Plan** had very little variation among the average values between clusters. This might signify that these attributes did not have much affect on the clustering process.

We can see that Cluster 1 had on average, higher **Income**, higher **Overage**, moderate **Leftover** minutes, higher **Handset Price**, higher number of **Over 15 Min Calls per Month**, and a higher number of "Leave" values for the **Leave** attribute. In contrast, Cluster 2 had on average, lower **Income**, lower **Overage**, lower **Leftover**, lower **Handset Price**, lower number of **Over 15 Min Calls per Month**, and a higher number of "Stay" values for the **Leave** attribute. Cluster 2 also has the most amount of points (6061 points), which shows that a large amount of the total points have lower than average values for their attributes.

Table 1: Cluster 1 (2212 points)

	College	Income	Overage	Leftover	House	Handset Price	Over 15 Min Calls per Month
Mean	0.509	128118	193.77	24.25	487948	641.66	18.86
Std	0.500	22094	42.80	26.79	250777	159.40	6.85
	Avg Call Duration		Reported Satis.	Reported Usage Lvl.		Consid. Change Plan	Leave
Mean	6.014		1.540	1.784		2.506	0.823
Std	4.408		1.623	1.494		1.317	0.381

Table 2: Cluster 2 (6061 points)

	College	Income	Overage	Leftover	House	Handset Price	Over 15 Min Calls per Month
Mean	0.509	56155	32.07	7.64	495074	263.20	2.60
Std	0.499	25797	39.57	9.76	251585	94.22	3.13
	Avg Call Duration		Reported Satis.	Reported Usage Lvl.		Consid. Change Plan	Leave
Mean	8.176		1.572	1.834		2.430	0.326
Std	3.947		1.625	1.517		1.338	0.468

Table 3: Cluster 3 (3716 points)

	College	Income	Overage	Leftover	House	Handset Price	Over 15 Min Calls per Month
Mean	0.487	66323	45.67	60.66	499972	305.00	3.52
Std	0.499	33610	56.78	18.82	254140	137.91	4.49
	Avg Call Duration		Reported Satis.	Reported Usage Lvl.		Consid. Change Plan	Leave
Mean	1.786		1.581	1.797		2.527	0.511
Std	1.424		1.625	1.514		1.301	0.499

Table 4: Cluster 4 (4106 points)

	College	Income	Overage	Leftover	House	Handset Price	Over 15 Min Calls per Month
Mean	0.491	56220	195.48	19.85	489171	263.13	19.29
Std	0.499	25241	40.70	23.47	252382	93.23	6.48
	Avg Call Duration		Reported Satis.	Reported Usage Lvl.		Consid. Change Plan	Leave
Mean	6.268		1.519	1.811		2.552	0.597
Std	4.321		1.627	1.513		1.314	0.490

Table 5: Cluster 5 (3905 points)

	College	Income	Overage	Leftover	House	Handset Price	Over 15 Min Calls per Month
Mean	0.513	129211	31.79	18.18	490828	656.55	2.59
Std	0.499	21348	37.56	21.58	252931	151.90	2.88
	Avg Call Duration		Reported Satis.	Reported Usage Lvl.		Consid. Change Plan	Leave
Mean	6.352		1.604	1.824		2.492	0.434
Std	4.251		1.643	1.509		1.329	0.495

6 Predictive Models

6.1 Decision Trees

The first type of predictive model I chose was the decision tree classifier. The decision tree model I used was `sklearn's DecisionTreeClassifier` function. The hyper-parameters I tuned were: Gini impurity to measure the quality of a split and a max tree depth of 10. This model was trained and tested with a 75/25 split on the dataset, with 75% of the dataset used for the training data. The decision tree classifier had on average 68% accuracy on the test data.

In an attempt to improve my testing accuracy, I decided to implement the random forest model, which is an ensemble learning method that improves upon the decision tree method by correcting their habit of over-fitting. Random forests do this by constructing a multitude of decision trees at training time and outputting the class that is the mode of the classes. The random forest model I used was `sklearn's RandomForestClassifier` function. Using the same data split as the decision tree classifier, my random forest model saw little improvement; the model had on average an accuracy of 70%.

The code for these models can be found in `../src/decisiontree.py`.

6.2 Neural Network

The next type of predictive model I chose was a simple feed-forward neural network. The model I used was a `keras Sequential` model. Like the decision tree model, I trained and validated the neural network using a 75/25 training and test split of the dataset. I created three different models, each with different hidden layer dimensions: **Figure 5:** 11x32x16x4x1, **Figure 6:** 11x64x32x16x4x1, and **Figure 7:** 11x128x64x32x16x8x4x1. The first layer in each model is 11 nodes for the 11 attributes, and the output layer is 1 node for the `Leave` attribute. All of my models trained on 20 epochs, used the binary cross-entropy loss function, and used the ReLU activation function on all layers except for the output layer, which used the sigmoid function.

The models had testing accuracies of 68.15%, 68.18%, and 67.78% for figures 5, 6, and 7, respectively.

The difference in accuracies is basically negligible, so we can see that the change in model dimensions doesn't affect the performance for this specific dataset. These also may not be the best hyper-parameters, so the model performance could be better.

The code for the neural network model can be found in `../src/neuralnet.py`.




Figure 5: Dimensions: 11x32x16x4x1, Test Accuracy: 68.15%




Figure 6: Dimensions: 11x64x32x16x4x1, Test Accuracy: 68.18%




Figure 7: Dimensions: 11x128x64x32x16x8x4x1, Test Accuracy: 67.78%

7 Analysis

Both the decision tree model and the neural network model achieved testing accuracies of about 68%. The random forest model, which was done to try and improve upon the decision tree model, achieved slightly better results, with a testing accuracy of 70%. All of these models were trained using 75% of the dataset, and tested on the remaining 25%. The data was shuffled before being split to try and mitigate any hidden patterns in the data ordering.

In general, the data was balanced; upon analysis of the values under the **Leave** attribute, there were 10,148 entries under the "STAY" class, and 9,853 entries under the "LEAVE" class. Since the data seems balanced, I don't have to worry about the impacts of un-balanced data on my model performance.

Some correlations I found between the data and churn is that